

Sequence Labelling

Natural Language Processing

Some slide content based on textbooks:

***Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition** by Daniel Jurafsky and James H. Martin*

ALICE was beginning to get very tired of sitting by her sister on the bank, and of having nothing to do: once or twice she had peeped into the book her sister was reading, but it had no pictures or conversations in it, "and what is the use of a book," thought Alice, "without pictures or conversations?" So she was considering in her own mind (as she could, for the hot day made her head heavy and stupid,) whether the pleasure of getting a daisy-chain would be worth the trouble of getting up and picking the daisies, when suddenly a white rabbit with pink eyes ran close to her. There was nothing so very remarkable in that; nor did Alice think it so very much out of the way to hear the Rabbit say to itself, "Oh dear! Oh dear! I shall be too late!" (whereupon she thought, over afterwards, it occurred to her that she ought to have wondered at this, but at the time she all seemed quite natural); but when the rabbit actually took a watch out of its waistcoat-pocket, and looked at it, and then hurried on, Alice started to her feet, for it flashed across her mind that she had never before seen a rabbit with either a waistcoat-pocket or a watch to take out of it, and

Minor Image source: <https://freemove.org/en/1574424884>

Lecture Contents:


- What is sequence labelling?
- Application: POS tagging
- Application: Named Entity Recognition
- How do sequence labellers work?
- Related applications for extracting information from text
 - Entity linking
 - Relation extraction

What is Sequence Labelling?

What is sequence labelling?


Classification task:

- input: **ordered sequence** of tokens: (x_1, x_2, \dots, x_n)
- output: single prediction for sequence: y



Sequence labelling task:

- input: **ordered sequence** of tokens: (x_1, x_2, \dots, x_n)
- output **sequence of predictions**: (y_1, y_2, \dots, y_n)



Note that:

- prediction for y_4 may depend not only on the sequence up to that point (x_1, x_2, x_3, x_4) , but also the subsequent sequence (x_5, \dots)
- dependencies exist across predicted sequence, so certain values y_4 may not make sense given values for (y_1, y_2, y_3) and y_5 .

Application:
part-of-speech (POS) tagging

What is POS tagging?

POS tagging:

- task of assigning to each token in a sequence:
- a **part-of-speech** label
- e.g.: **PRON** (pronoun), **VERB**, **DET** (determiner), **NOUN**, etc.

Why label parts-of-speech?

- useful for **developing features** for certain tasks
 - e.g. authorship attribution, particularly if only small amount of training data is available
- useful to **reduce ambiguity** in bag-of-words representation
 - some terms have different meaning depending on context “to book” vs “a book”
 - so append POS tag to each word occurrence: book_VERB vs book_NOUN
- useful as **initial step** for other NLP tasks or performing linguistic analysis
 - required for syntactic parsing
 - useful for text-to-speech
 - pronouncing “lead group” vs “lead weight” or “to object” vs “an object”
 - studying linguistic change like creation of new words, or meaning shift

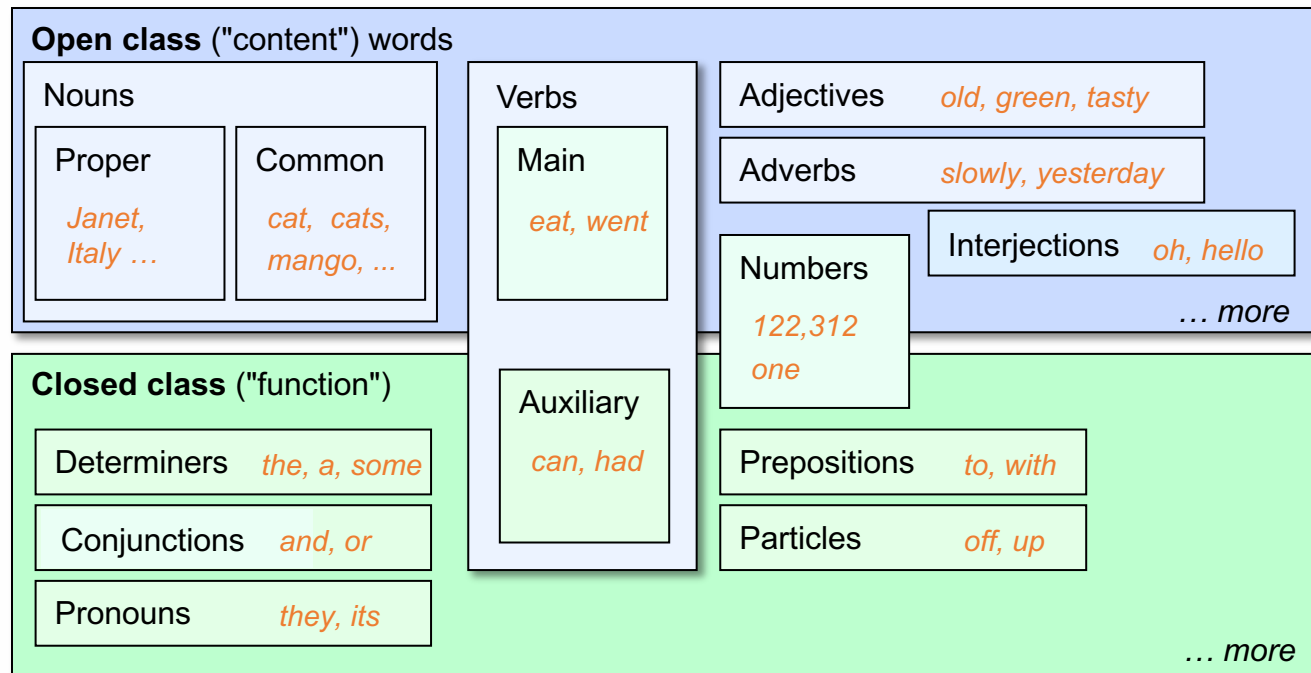
[(He, 'PRON'),
(thought, 'VERB'),
(he, 'PRON'),
(saw, 'VERB'),
(an, 'DET'),
(elephant, 'NOUN'),
(riding, 'VERB'),
(a, 'DET'),
(bicycle, 'NOUN'),
(on, 'ADP'),
(the, 'DET'),
(freeway, 'NOUN')]

Parts of Speech classes

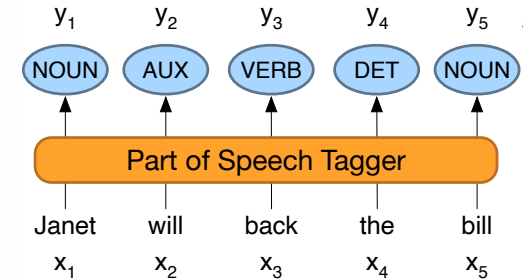
Word classes have been around a **long time**:

- Back in the 1st century BCE, Dionysius Thrax of Alexandria defined:
nouns, verbs, pronouns, prepositions, adverbs, conjunctions, participles, articles

Modern grammar divides word classes into open and closed:



Parts of Speech tagging



Map sequence of words x_1, \dots, x_n to sequence of POS tags y_1, \dots, y_n

- set of tags:

	Tag	Description	Example
Open Class	ADJ	Adjective: noun modifiers describing properties	<i>red, young, awesome</i>
	ADV	Adverb: verb modifiers of time, place, manner	<i>very, slowly, home, yesterday</i>
	NOUN	words for persons, places, things, etc.	<i>algorithm, cat, mango, beauty</i>
	VERB	words for actions and processes	<i>draw, provide, go</i>
	PROPN	Proper noun: name of a person, organization, place, etc..	<i>Regina, IBM, Colorado</i>
	INTJ	Interjection: exclamation, greeting, yes/no response, etc.	<i>oh, um, yes, hello</i>
Closed Class Words	ADP	Adposition (Preposition/Postposition): marks a noun's spacial, temporal, or other relation	<i>in, on, by under</i>
	AUX	Auxiliary: helping verb marking tense, aspect, mood, etc.,	<i>can, may, should, are</i>
	CCONJ	Coordinating Conjunction: joins two phrases/clauses	<i>and, or, but</i>
	DET	Determiner: marks noun phrase properties	<i>a, an, the, this</i>
	NUM	Numeral	<i>one, two, first, second</i>
	PART	Particle: a preposition-like form used together with a verb	<i>up, down, on, off, in, out, at, by</i>
	PRON	Pronoun: a shorthand for referring to an entity or event	<i>she, who, I, others</i>
Other	SCONJ	Subordinating Conjunction: joins a main clause with a subordinate clause such as a sentential complement	<i>that, which</i>
	PUNCT	Punctuation	<i>; , ()</i>
	SYM	Symbols like \$ or emoji	<i>\$, %</i>
	X	Other	<i>asdf, qwfg</i>

- example sentences:

- There/**PRO** were/**VERB** 70/**NUM** children/**NOUN** there/**ADV** ./**PUNC**
- Preliminary/**ADJ** findings/**NOUN** were/**AUX** reported/**VERB** in/**ADP** today/**NOUN** 's/**PART** New/**PROPN** England/**PROPN** Journal/**PROPN** of/**ADP** Medicine/**PROPN**

Is POS tagging difficult?

Approximately 85% of vocabulary terms in English are unambiguous

- *Janet* is always PROP, *hesitantly* is always ADV

But ambiguous vocab terms are very common

- so ~60% of tokens are ambiguous

Example: word **back** could have 5 different POS tags:

- *earnings growth took a back/ADJ seat*
- *a small building in the back/NOUN*
- *a clear majority of senators back/VERB the bill*
- *enable the country to buy back/PART debt*
- *I was twenty-one back/ADV then*

Accuracy of POS tagging is about 97%

- changed little in last 10+ years: HMMs, CRFs, and BERT perform similarly
- similar to human accuracy
- baseline (label each word with its most frequent tag) performance already 92%

Features used for POS tagging

Consider the example:

- Janet *will* back the *bill*
AUX/NOUN/VERB? NOUN/VERB?

Sources of evidence for determining the POS tags:

- Prior probabilities of word/tag
 - "will" is usually an AUX
- Identity of neighboring words
 - "the" means the next word is probably not a verb
- Morphology and wordshape:
 - Prefixes *unable:* un- → ADJ
 - Suffixes *importantly:* -ly → ADJ
 - Capitalization *Janet:* CAP → PROP

Application:
named-entity-recognition (NER)

What is entity recognition?

Named-Entity Recognition (NER):

- task of identifying entities that are mentioned in a text
- can be treated as a sequence labelling task
- often a first step in extracting knowledge from text

"Have you heard of an associate professor from
the Politecnico di Milano called Mark Carman?"

Institution Person

Named Entity Recognition

Named entity = object in real world

- most common tags:
 - **PER** (Person): e.g. "Marie Curie"
 - **LOC** (Location): e.g. "Lake Michigan"
 - **ORG** (Organization): e.g. "Stanford University"
 - **GPE** (Geo-Political Entity): e.g. "Boulder, Colorado"
- often multi-word phrases
- term also extended to things that aren't entities:
 - dates, times, prices

Difference between a GPE and a LOC?

- **GPE**: geopolitical entities, e.g. everything with a governing body like cities and countries. Examples: "Germany", "Buenos Aires".
- **LOC**: everything else that's a physical location or area, like "Kalahari Desert" or "Silicon Valley"

Source: <https://support.prodi.gov.it/ner-annotation-scheme-gpe-vs-loc/2913>

NER task: find spans of text that constitute proper names

- tagging the type of the entity:

Citing high fuel prices, [ORG United Airlines] said [TIME Friday] it has increased fares by [MONEY \$6] per round trip on flights to some cities also served by lower-cost carriers. [ORG American Airlines], a unit of [ORG AMR Corp.], immediately matched the move, spokesman [PER Tim Wagner] said. [ORG United], a unit of [ORG UAL Corp.], said the increase took effect [TIME Thursday] and applies to most routes where it competes against discount carriers, such as [LOC Chicago] to [LOC Dallas] and [LOC Denver] to [LOC San Francisco].

Why NER?

Traditionally perform NER for:

- **Sentiment analysis:** identify sentiment towards particular company or person?
- **Information extraction:** extracting facts about entities from text
- **Question answering:** answer questions about an entity?
- **De-identification:** remove references to individual from text to protect privacy

NER is hard because of:

- 1) segmentation: in POS tagging each word gets one tag, while in NER have to find and segment entities
- 2) type ambiguity: same word/phrase could have many types depending on the context

[PER Washington] was born into slavery on the farm of James Burroughs.
[ORG Washington] went up 2 games to 1 in the four-game series.
Blair arrived in [LOC Washington] for what may well be his last state visit.
In June, [GPE Washington] passed a primary seatbelt law.

Begin-Inside-Outside (BIO) Tagging

NER finds phrases in the text referring to named entities:

- *[PER Jane Villanueva] of [ORG United] , a unit of [ORG United Airlines Holding] , said the fare applies to the [LOC Chicago] route.*

How can we turn NER into sequence labeling problem (with one label per token)?

- use begin/inside/outside tags:
 - **B**: token that **begins** a span
 - **I**: tokens **inside** a span
 - **O**: tokens **outside** of any span


Words	BIO Label
Jane	B-PER
Villanueva	I-PER
of	O
United	B-ORG
Airlines	I-ORG
Holding	I-ORG
discussed	O
the	O
Chicago	B-LOC
route	O
.	O

Sequence Labelling: approaches

How do sequence labellers work?

Traditional methods for sequence labelling make use of

- Hidden Markov Models (HMMs) = Naïve Bayes applied to sequences



- Conditional Random Fields (CRFs) = Logistic Regression applied to sequences

Recent methods make use of

- Recurrent Neural Networks (RNNs) to improve performance

Recurrent Neural Networks (RNNs)

RNNs build upon word embeddings

- by aggregating information along sequences 🙌

Provide general mechanism for combining:

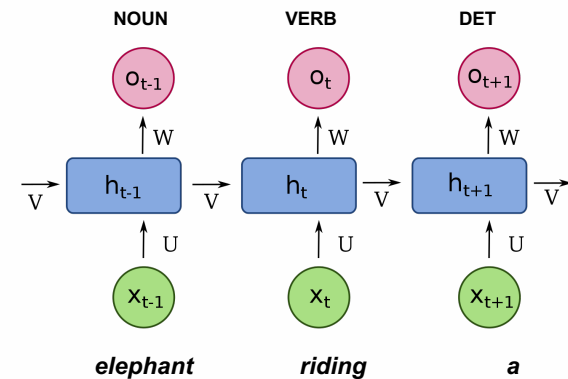
- context from previous words
- with **embedding** of current word

Implemented as NN which:

- takes 2 input vectors: (current input, previous state)
- produce 2 output vectors: (current output, updated state)

By updating an internal state, RNN is able to:

- process arbitrarily long inputs
- produce output prediction at each input position



Source: https://en.wikipedia.org/wiki/Recurrent_neural_network

Aside: RNNs and word order

Word order is very important for interpreting the meaning of text

- and interpreting the meaning is important **for classifying it**

For example consider the **meaning** of the following phrases:

- There's a white rat in the house ...
- There's a rat in the White House ...

Negation provides another important example of word order:

- I am happy about ...
- I am not happy about ...
- I'd be lying if I said I was not happy about ...
- I would not be lying if I said that I was not happy about ...

N-grams can be used to capture word order

- **but** we can never make them long enough



Source: https://commons.wikimedia.org/wiki/File:White_rat_on_table.jpg



Source: https://commons.wikimedia.org/wiki/File:White_House_DC.JPG

Long Short-Term Memory (LSTM)

Clever implementation of RNN that is able to


- learn **contexts** and **long range dependencies**

Does this by using a **gating mechanism**

- **passes through information by default**
- unless new information is added to state
- or deleted from it (forgotten)

LSTM learns when & what information to

- **remember**, **forget**, and **output** at each timestep



Images source:
Understanding LSTM Networks by Christopher Olah
<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Aside: LSTMs and handling context

LSTMs can be stacked on top of each other

- have uncanny ability to handle nested contexts
- useful for natural language:
 - for example, complete sentences with: *he/she/his/her*
My mother was taking on the phone to ___ friend Jim.
Jim said that ___ favourite game was confusing ___ students.
Replying, ___ said that ___ should find a better hobby.
 - gender of subject changes for each subsequent sentence
 - another example, this time with negation, complete with: *friendly/self-absorbed*
I get along well with her brother. He's always ___
I can not get along well with her brother. He's always ___
I can not help but get along well with her brother. He's always ___
 - LSTMs are able to switch between sentence and negation contexts

For $\mathbb{G}_{n+1, \dots, n}$ where $\mathcal{L}_{n+1} = 0$, hence we can find a closed subset \mathcal{H} in \mathcal{H} and any sets \mathcal{F} on X , U is a closed immersion of S , then $U \rightarrow T$ is a separated algebraic space.

Proof. Proof of (1). It also start we get

$$S = \text{Spec}(R) = U \times_X U \times_X U$$

and the comparably in the fibre product covering we have to prove the lemma generated by $\prod \mathcal{F}_i \times_U U \rightarrow V$. Consider the maps M along the set of points Sch_{aff} and $U \rightarrow U$ is the fibre category of S in U in Section 77 and the fact that any U affine, see Morphisms, Lemma 77. Hence we obtain a scheme S and any open subset $W \subset U$ in $\text{Sch}(G)$ such that $\text{Spec}(R') \rightarrow S$ is smooth or an

$$U = \bigcup U_i \times_S U_i$$

which has a nonzero morphism we may assume that f_i is of finite presentation over S . We claim that $\mathcal{O}_{X, s}$ is a scheme where $x, x', s'' \in S'$ such that $\mathcal{O}_{X, s} \rightarrow \mathcal{O}_{X', s'}$ is separated. By Algebra, Lemma 77 we can define a map of complexes $\text{GL}_S(x'/S')$ and we win.

To prove study we see that $\mathcal{F}_i|_U$ is a covering of X' , and \mathcal{T}_i is an object of $\mathcal{F}_{X/S}$ for $i > 0$ and \mathcal{F}_0 exists and let \mathcal{F}_i be a presheaf of \mathcal{O}_X -modules on \mathcal{C} as a \mathcal{F} -module. In particular $\mathcal{F} = U/\mathcal{F}$ we have to show that

$$\tilde{M}^* = \mathbb{Z}^* \otimes_{\text{Spec}(R)} \mathcal{O}_{S, s} \rightarrow \Gamma_S^{-1}(\mathcal{F})$$

is a unique morphism of algebraic stacks. Note that

$$\text{Arrows} = (\text{Sch}/S)_{\text{aff}}^{\text{op}} / (\text{Sch}/S)_{\text{aff}}$$

and

$$V = \Gamma(S, \mathcal{O}) \rightarrow \Gamma(U, \text{Spec}(A))$$

is an open subset of X . Thus U is affine. This is a continuous map of X is the inverse, the groupoid scheme S .

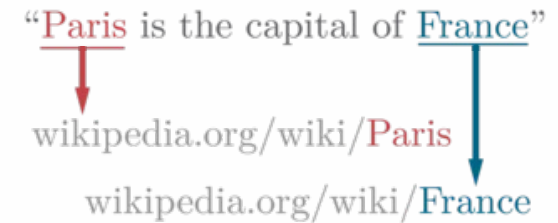
Proof. See discussion of sheaves of sets.

The result for prove any open covering follows from the loss of Example 77. It may replace S by $X_{\text{space}, \text{finite}}$ which gives an open subspace of X and T equal to S_{space} , see Descent, Lemma 77. Namely, by Lemma 77 we see that R is geometrically regular over S . □

PDF compiled from automatically generated LaTeX using multi-layer LSTM by Andrej Karpathy
<http://karpathy.github.io/2015/05/21/mn-effectiveness/>

Entity Linkage

What is entity linkage?



Source: https://commons.wikimedia.org/wiki/File:Entity_Linking_-_Short_Example.png

Determining that a named-entity has been mentioned in text

- often **only the first part** of problem
- second part: determine **which real-word entity** was referred to
- not as easy as it sounds!

Linkage techniques make use of:

- relative importance of entities
- context within text (other entities present)

Ontology/Knowledge Base

- generally **Wikipedia/DBPedia** is used
 - but many individuals/objects have no Wikipedia page
 - so custom custom ontologies can be used

I grew up in a small town just out of Paris.
Currently driving from Dallas to Paris.
Paris Hilton was photographed leaving the Paris
Hilton.

Just had my photo taken with Michael Jordan!!
Just had my photo taken with Michael Jordan at EMNLP!!

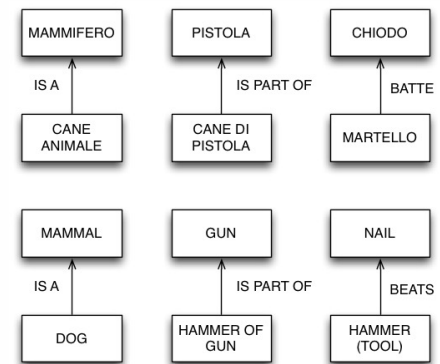
Aside: Taxonomies and Ontologies

What are taxonomies and ontologies?

Taxonomy = hierarchy of concepts (e.g. types of products with is-a or part-of relationships)

Ontology = formal definition of concepts belonging to a domain

- Abstract definition of concepts that does not depend on the language




Most ontologies are composed of:


- Classes (e.g. **wine, winery**): A set of elements; a type
- Individuals (e.g. **champagne**): An element/object
- Attributes (e.g. **price**): property with primitive data type (e.g. **string/integer**) allowing for restrictions on values (e.g., “>0”)
- Relationships (e.g. **winery produces wine**): characterization of relationships among classes or individuals
- Logical rules, e.g.:
$$\text{hasParent} (?x1, ?x2) \wedge \text{hasBrother} (?x2, ?x3) \rightarrow \text{hasUncle} (?x1, ?x3)$$

Ontologies as graphs

The relationships between concepts in an ontology/knowledge base form a graph:



- Ontologies/knowledge bases can be used to represent information (called “facts”) contained in sentences
 - e.g. for the sentence “*I have a car*”



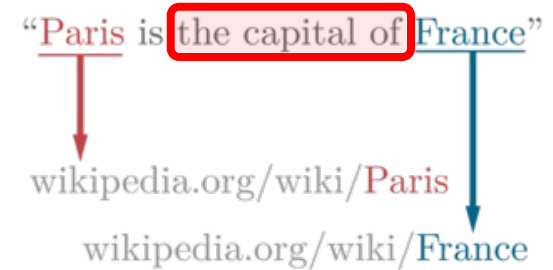
Web Ontologies and Knowledge Bases

OWL (Web Ontology Language)

- based on RDF (triple: *subject predicate object*), i.e. Description Logic
- uses SPARQL query language to allow inference over KB such as DBPedia
- KBs have **open world** semantics:
 - any statement that is not known to be true is unknown
 - as opposed to closed world assumption used in SQL:
 - any statement that is not known to be true is false (*negation as failure*)
 - example if KB contains propositions:
“Giovanni is an architect”, “Giovanni is not a physicist”
 - query: “Is Giovanni an engineer?”
open world answer: *unknown*, closed world answer: *no* (proposition not in KB)
 - query: “Is Giovanni a physicist?”
open world answer: *no* (negated proposition found), closed world answer: *no* (proposition not in KB)

Relation Extraction

What is relation extraction?



Source: https://commons.wikimedia.org/wiki/File:Entity_Linking_-_Short_Example.png

Once entity mentions have been linked to unique entities

- **relationships between entities** can be mined
- and used to populate a knowledge graph / knowledge base

Handled as a problem of predicting **missing links** in a graph

- entity embeddings can be leveraged for this purpose
 - since translations in space naturally encode relationships
 - ongoing research topic, see e.g. <https://arxiv.org/pdf/2002.00388.pdf>


Relation: capitalOf	
city	country
London	UK
Rome	Italy
Paris	France
Canberra	Australia
Belgrade	a Serbia

Parse Trees

Parse Tree

Parse Trees result from applying a **formal grammar** to understand how a sentence was generated

- used to be very popular in NLP



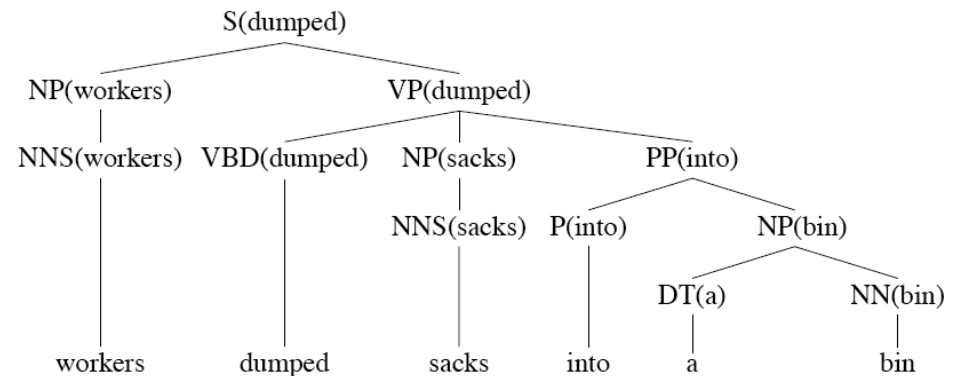
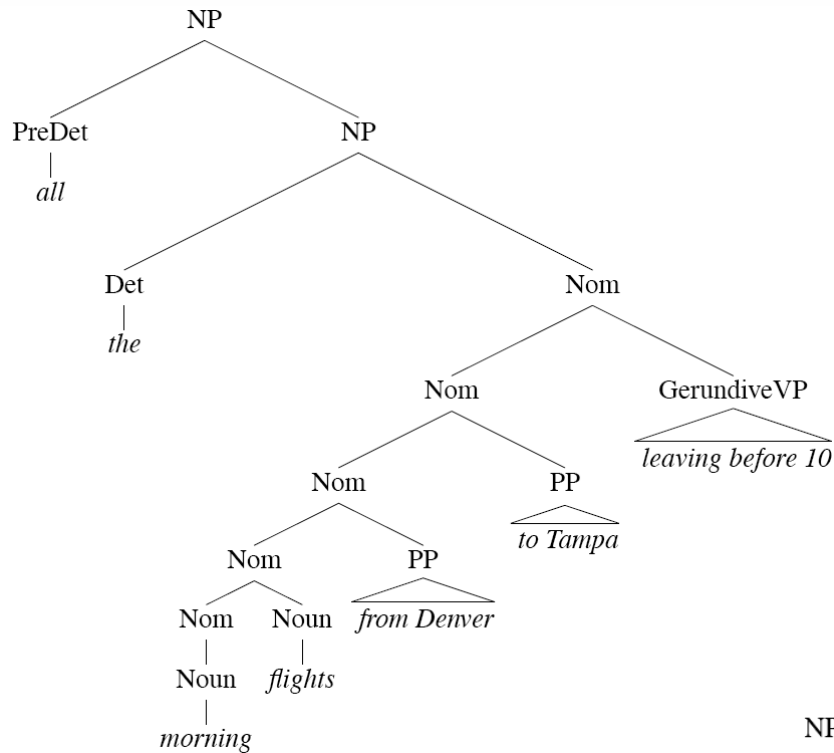
Penn Treebank

Treebank

- corpora in which sentences are paired with a parse tree
- Most famous is the Wall Street Journal section of the Penn TreeBank.
 - One million words from Wall Street Journal.

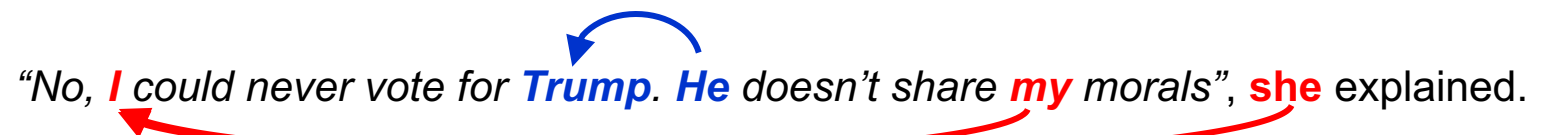
```
( (S ( ' ' ' ' )
  (S-TPC-2
    (NP-SBJ-1 (PRP We) )
    (VP (MD would)
      (VP (VB have)
        (S
          (NP-SBJ (-NONE- *-1) )
          (VP (TO to)
            (VP (VB wait)
              (SBAR-TMP (IN until)
                (S
                  (NP-SBJ (PRP we) )
                  (VP (VBP have)
                    (VP (VBN collected)
                      (PP-CLR (IN on)
                        (NP (DT those)(NNS assets))))))))))
          ( , , ) ( ' ' ' ' )
          (NP-SBJ (PRP he) )
          (VP (VBD said)
            (S (-NONE- *T*-2) ))
          ( . . ) ))
```


Examples of Parse Trees:



Co-reference resolution

“No, *I* could never vote for *Trump*. *He* doesn’t share *my* morals”, *she* explained.



The diagram illustrates co-reference resolution in the sentence: “No, *I* could never vote for *Trump*. *He* doesn’t share *my* morals”, *she* explained. A blue curved arrow points from the pronoun *He* to the name *Trump*. A red curved arrow points from the pronoun *she* to the pronoun *I*. Another red curved arrow points from the pronoun *my* to the pronoun *I*.

Co-reference Resolution

Problem of determining who or what is being referenced across (or sometimes within) sentences:

***John** went to **Bill's car dealership** to check out **an Acura Integra**.
He looked at **it** for half an hour*

- In the second sentence, who is being referred to by the word **he**?
 - *He* refers to *John*
- In the second sentence, what is **it**?
 - Is it *Bill's car dealership*?
 - Or *An Acura Integra*?

Co-reference Resolution (cont.)

- Sometimes pronoun comes after referent (anaphora):
John went to the dealership to see a car that he was interested to purchase
- Sometimes pronoun comes before referent (cataphora):
Before he bought it, John checked over the Integra very carefully

Why resolve co-references to entities from earlier/later in the text?

- In order to be able to understand what is being said about those entities when pronouns are being used
- So needed for chatbots and information extraction ...

Types of Reference Phenomena

Pronouns:

I saw no less than 6 Acura Integras today. They are the coolest cars.

Demonstratives

I bought an Integra yesterday, similar to the one I bought five years ago. That one was nice, but I like this one even more

- *This* and *that* often refer metaphorically to time

A non-pronominal anaphora

I saw no less than 6 Acura Integra today. I want one

- ... one (of them)

Inferable

I almost bought an Acura Integra today, but the engine seemed noisy.

- The engine of...? Easy to infer: the Acura Integra

Usually resolve references to entities, but other things can be referenced in text:

According to John, Bob bought Sue an Integra, and Sue bought Fred a Legend

- But *that* turned out to be a lie (a speech act)
- But *that* was false (proposition)
- *That* struck me as a funny way to describe the situation (manner of description)
- *That* caused Sue to become rather poor (event)

Conclusions

Conclusions

TODO